

# Developing with AI Agent Swarms

Agentic software development  
with claude-code and  
claude-flow



# Dive Deep Into Multi-Agent Systems!

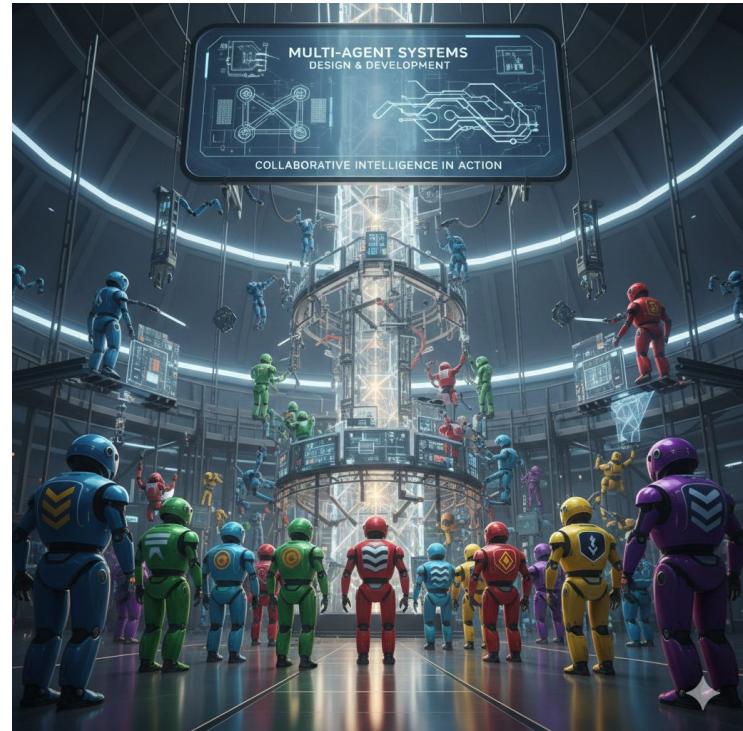
From Simple Workflows to Self-Organizing Swarms

## Understand multi-agent systems

- What can you do with them?
- What are the architectures involved?
- How do you make the right decisions?

Complete GitHub repository with slides, script, and demo instructions available at:

<https://github.com/jamesurquhart/swarmclass>



# Meet your Instructor

James Urquhart, Field CTO, KamiwazaAI

35 years of distributed systems design, development, deployment, and operations experience.

Senior technical leadership roles at AWS, VMware, Pivotal, Cisco, Dell, Enstratius, and Forte Software

Author of *Flow Architectures: The Future of Streaming and Event-Driven Integration* (O'Reilly, 2021)

Author of The Wisdom of Clouds from 2006-2016.  
Consistently voted one of the top three most influential blogs on cloud computing.

Currently Field CTO for KamiwazaAI, delivering AI data access, governance, and security to the enterprise



image: James Urquhart



# Course schedule

## Section 1:

### **Agentic AI and Swarm Basics (60 min)**

- Concepts, terminology, six architectures
- Demo: Sequential Pipeline (Research & Write)

## Section 2:

### **Creating Advanced Multi-Agent Systems (60 min)**

- Deep dives: Hierarchical, Debate, Committee, Routing
- Demo: Multi-Agent Code Review

## Section 3:

### **Applying Multi-Agent Patterns (60 min)**

- Decision framework, emergent coordination, best practices
- Demo: Self-Organizing Swarm walkthrough

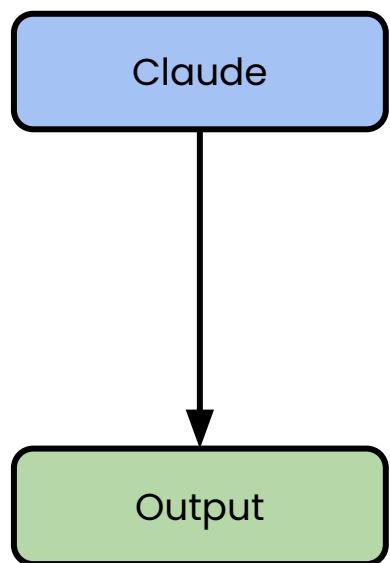
# Course Objectives

**By the end of this course, you will:**

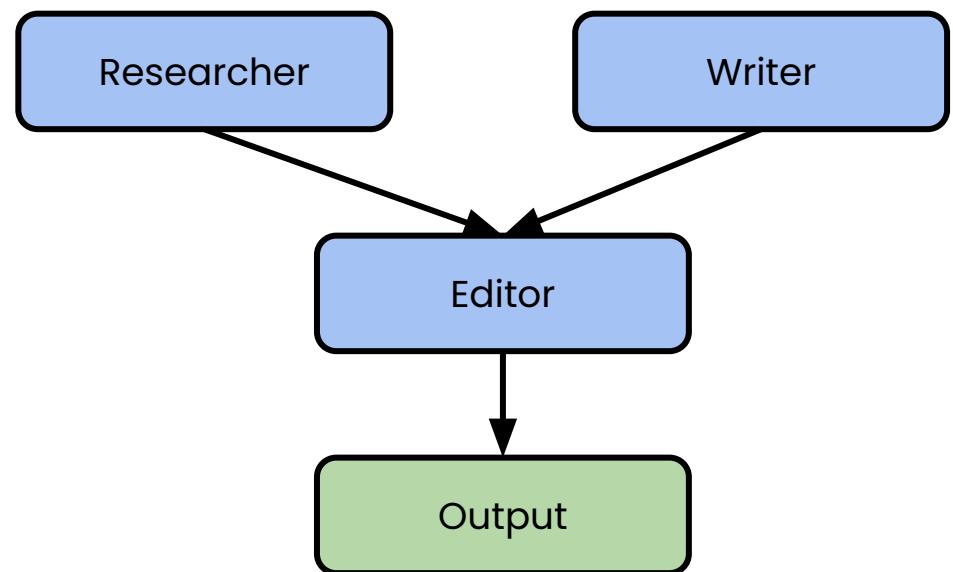
- Build multi-agent systems and see swarm coordination in action with claude-flow
- Know six architectural patterns and when to choose swarms over other approaches
- Understand how swarms differ from structured multi-agent orchestration

# The Shift

## Single Agent



## Multi-Agent



# Definition

**Multi-agent AI** is a computational paradigm where two or more autonomous AI **agents**—each with distinct roles, capabilities, or knowledge domains—coordinate, collaborate, or compete to achieve individual or collective goals through structured **interaction protocols**.

# Why Multi-Agent Systems?

## Single Agent Limitations

Context window constraints



## Multi-Agent Solutions

Decompose across agents

Lack of specialized depth



Each agent optimized for one task

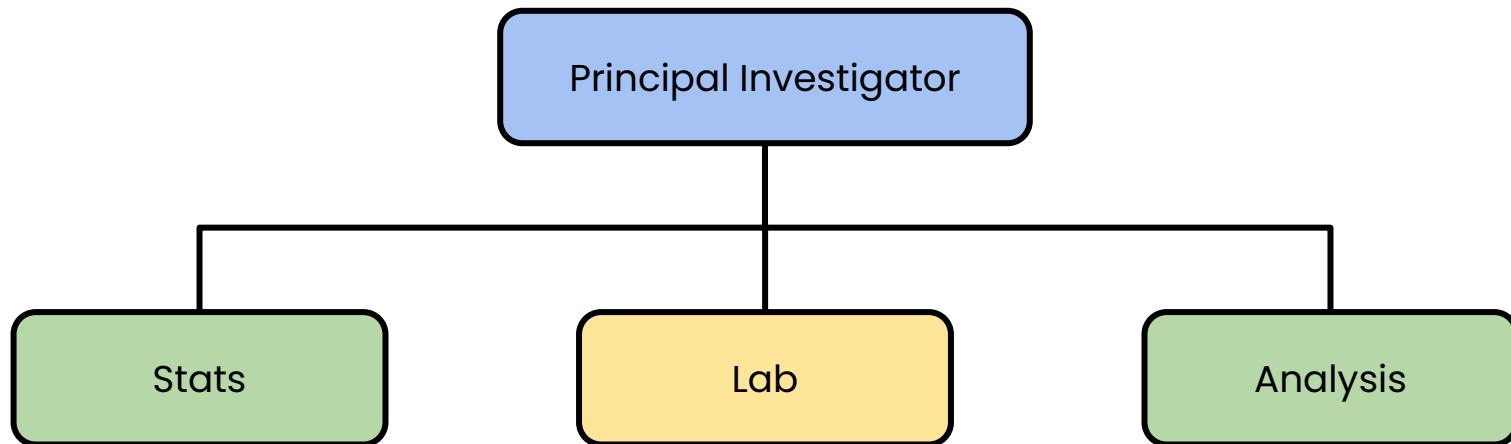
Sequential bottlenecks



Parallel execution

# The Human Analogy

How effective human teams work:



- Specialists coordinate
- They debate and verify each other
- Outcomes none could achieve alone



# Real-World Deployments

## Financial Services

Multi-agent fraud detections

*95%+ accuracy, reduced  
false positives*

## Insurance

7-agent claims processing

*80% faster processing time*

## Warehouses

Amazon Kiva robot swarms

*Hundreds of agents  
coordinating in real-time*

## AI Products

ChatGPT (Operator + Deep  
Research + Code)

*Claude Code (subagent  
spawning)*

## Edge/On-Prem

Local open source models  
(Llama, Mistral)

*Privacy, performance, and  
air-gapped deployments*



# Core Definitions

## Agent

Autonomous AI with specific role, tools, and decision-making capability

## Subagent

An agent spawned by another agent to handle a subtask

## Orchestration

Coordinating multiple agents toward a shared goal

## Context Window

The information an agent can process at once; each agent has its own

## Handoff

Passing work from one agent to another

## Shared Memory

A common data store agents use to communicate state



# Core Definitions

## MCP

Model Context Protocol—open standard for connecting agents to external tools/data

## Swarm

Multiple agents working in parallel with emergent coordination (not centrally directed)

# Framework Examples

Framework	Philosophy	Best For
<b>Claude Code</b>	Terminal-based agents	Development workflows with tool access
<b>Claude Flow</b>	Swarm orchestration,	Parallel agent shared memory coordination
<b>LangGraph</b>	Graph-based stateful workflows	Complex orchestration with precise control
<b>CrewAI</b>	Role-based teams	Rapid prototyping ("researcher", "writer")



## Emerging Standard:

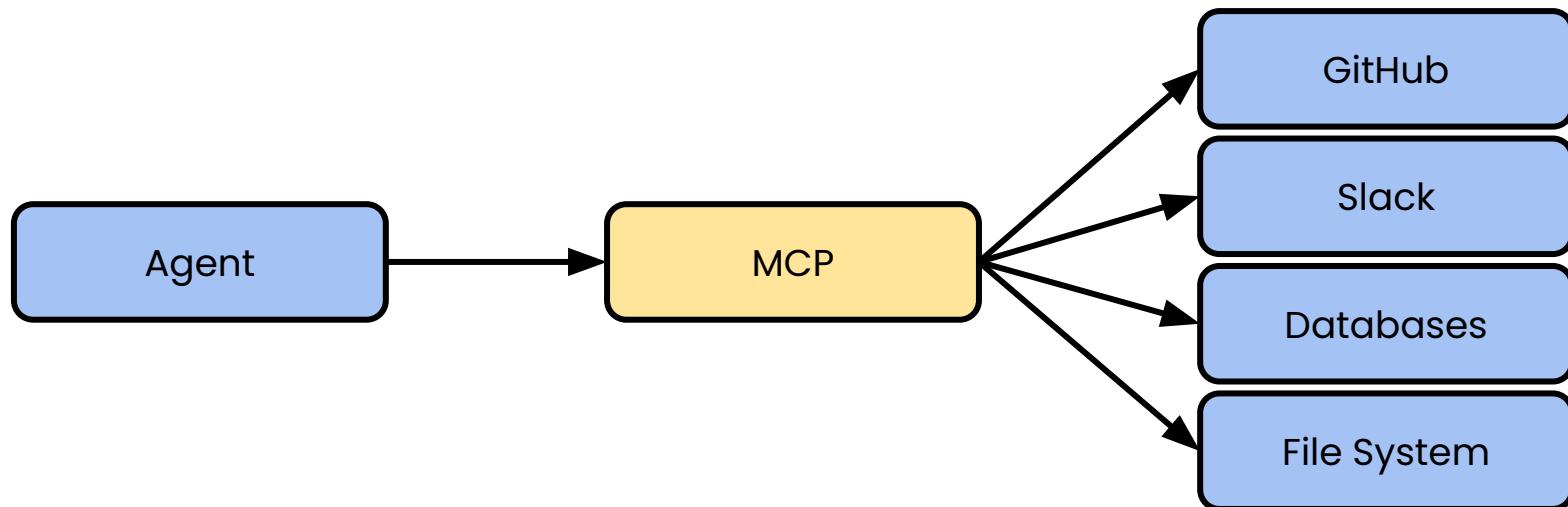
Agentic AI Foundation (Linux Foundation)

→ MCP, goose, AGENTS.md

→ Backed by Anthropic, OpenAI, Google, Microsoft, AWS

# MCP – The Connector

## Model Context Protocol



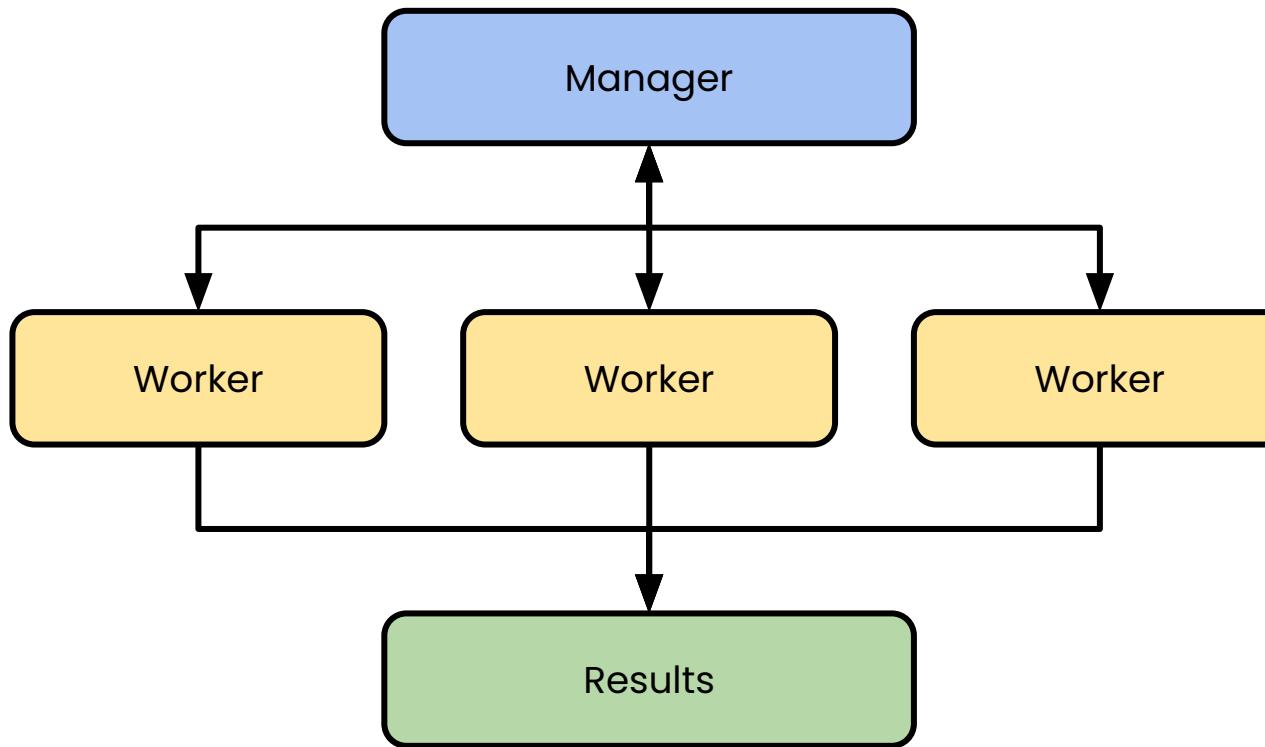
Originally Anthropic • Now AAIF/Linux Foundation

10,000+ published MCP servers

# Six Multi-Agent Architectures

Pattern	Description
Hierarchical	Manager delegates to workers
Debate	Agents argue, judge decides
Committee	Parallel processing, aggregated results
Pipeline	Sequential handoffs
Routing	Dynamic dispatch to specialists
Network/Swarm	Emergent coordination

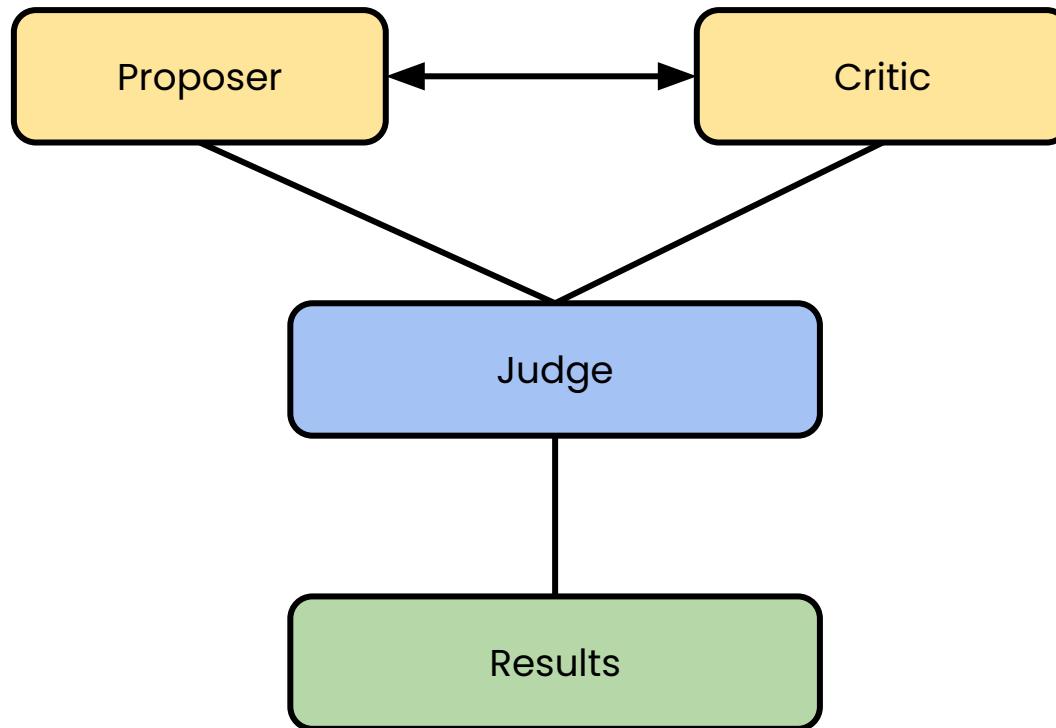
# Hierarchical



**Control:**  
Centralized

**Best for:**  
Decomposable  
tasks

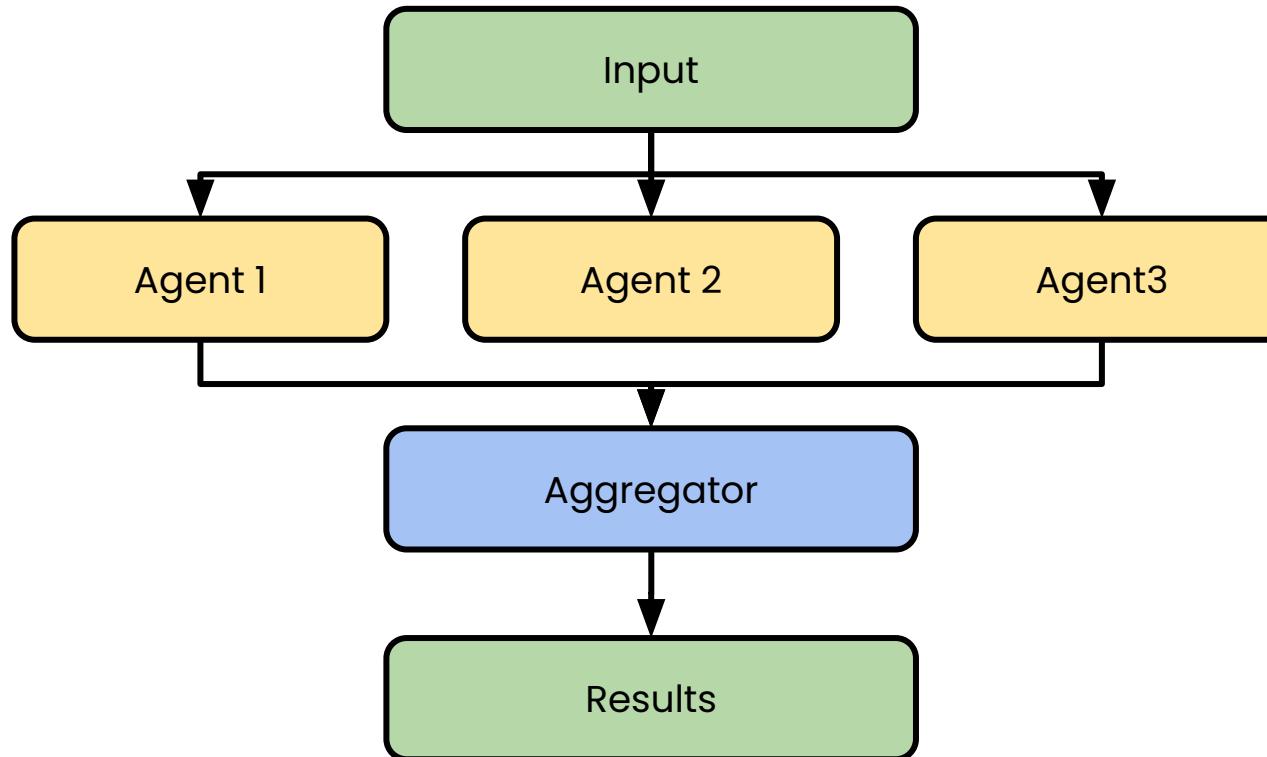
# Debate (Adversarial)



**Control:**  
Structured conflict

**Best for:**  
High-stakes decisions

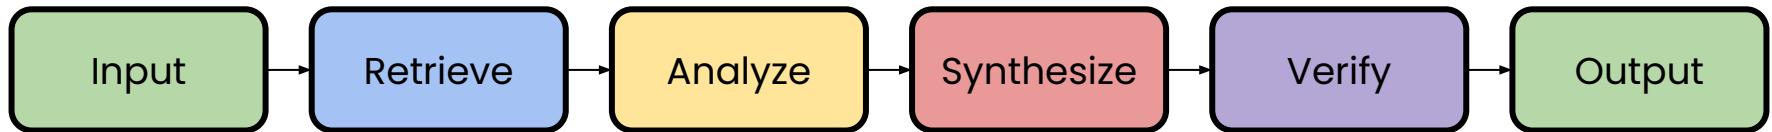
# Committee



**Control:**  
Parallel

**Best for:**  
Reliability-critical  
tasks

# Pipeline (Sequential)

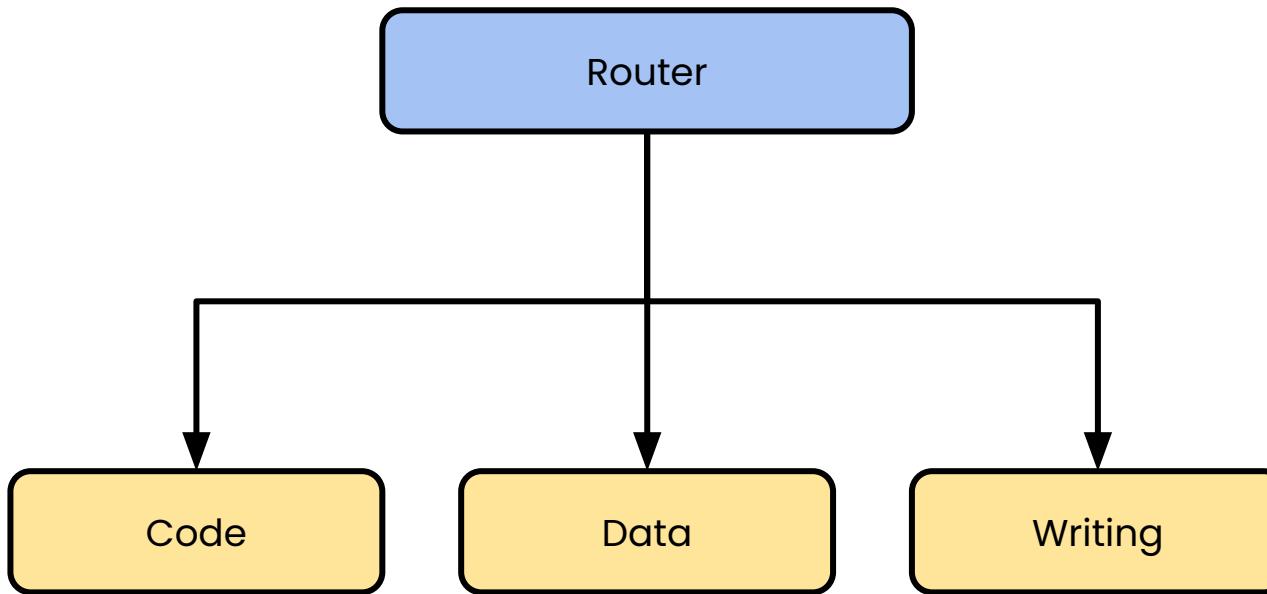


**Control:**  
Sequential

**Best for:** Staged transformations

(Our first demo uses this pattern)

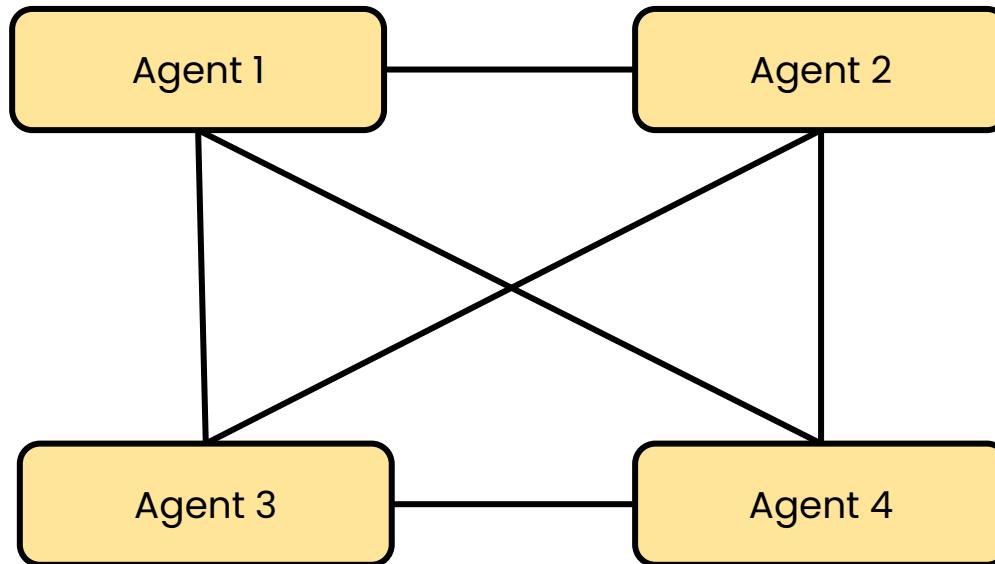
# Dynamic Routing



**Control:**  
Adaptive

**Best for:**  
Heterogeneous  
queries

# Network/Swarm



**Control:**  
Emergent

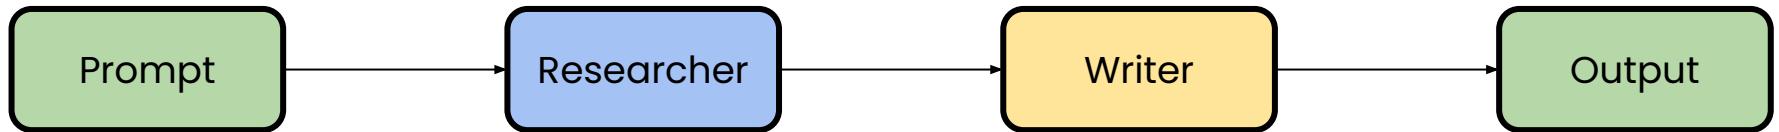
**Best for:**  
Novel  
problem-solving

# Architectures at a Glance

Architecture	Control	Best For
Hierarchical	Centralized	Decomposable tasks
Debate	Structured	Nuanced decisions
Committee	Parallel	Performance and reliability
Pipeline	Sequential	Workflows
Routing	Selective	Specialization and complex activities
Swarm/Hive	Emergent	Novel problem-solving

# Demo

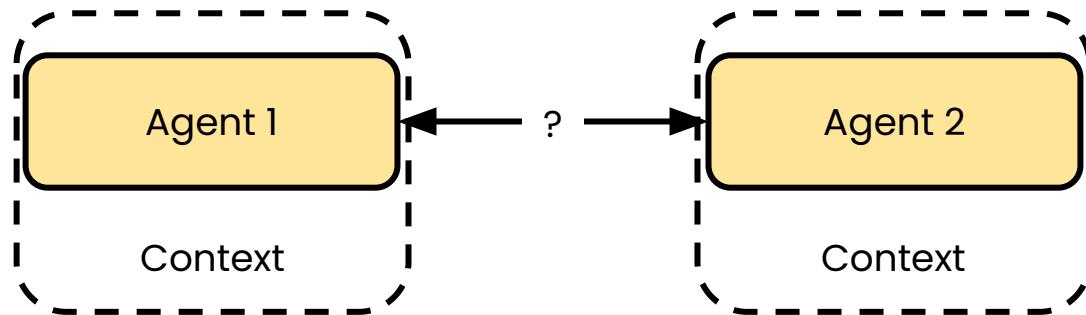
## Research and Write Pipeline



Claude-code

Anthropic

# Sharing Context and Signals



Each agent has its own context window.

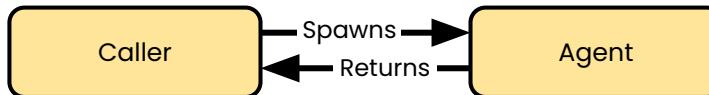
How do they share information?

# Four Handoff Patterns

## 1. File-Based (our demo)

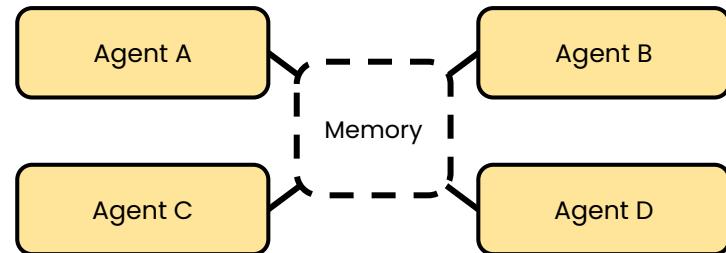


## 2. Return Value

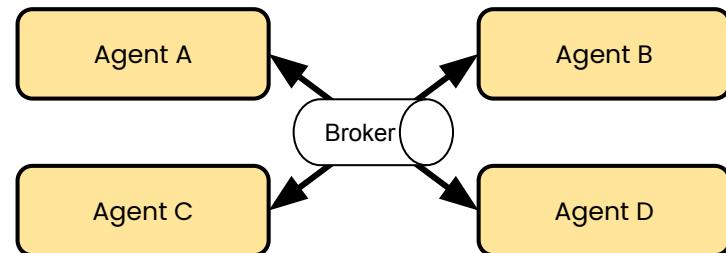


Caller could be an parent agent, an orchestrator, or the main session

## 3. Shared Memory



## 4. Message Passing



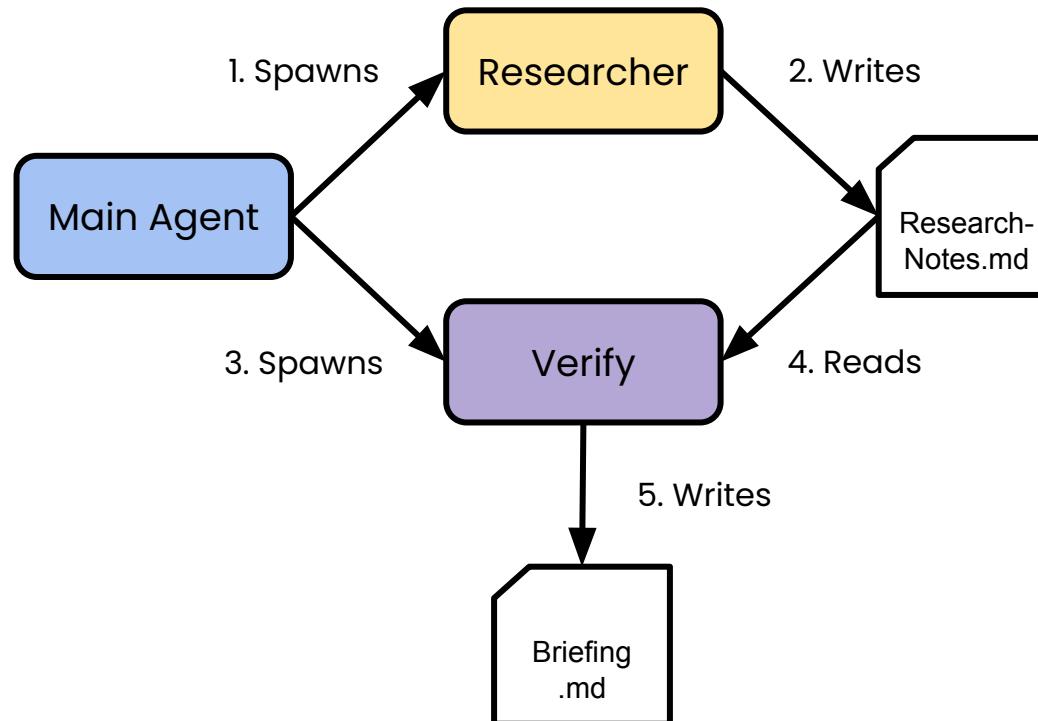
Requires external infrastructure—not native to Claude code

<b>Pattern</b>	<b>Coupling</b>	<b>Debugability</b>	<b>Parallel</b>	<b>Use When</b>
File-Based	Loose	High	No	Clear phases
Return Value	Medium	High	No	Call-Response
Shared Memory	Loose	Med	Yes	Swarms/Parallel
Message Passing	Tight	Low	Yes	Real-time*

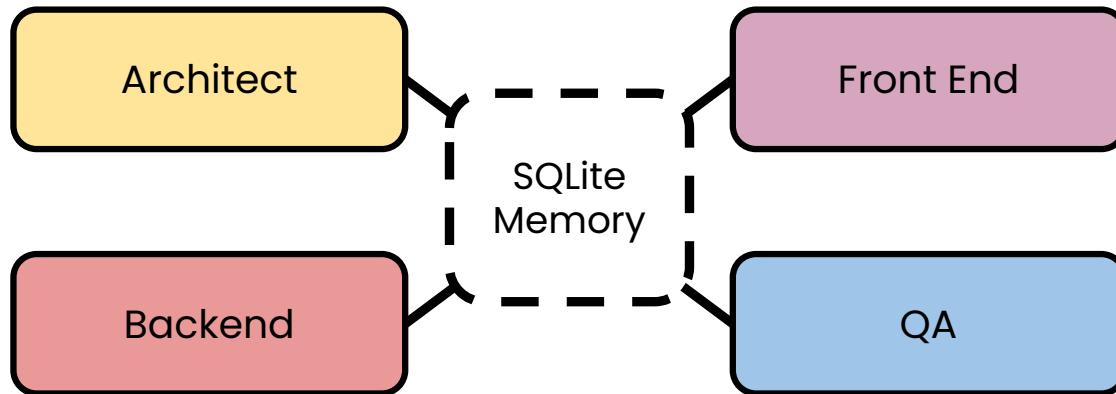
\* Requires external infrastructure



# Our Demo's Handoff



Coming in Section 2: Shared Memory Swarms



Agents discover and build on each other's work  
No explicit handoffs required

# Pulse Check

These are a way to quickly check in with attendees. Ask them a simple yes or no question, and the platform will prompt them to press “thumbs up” or “thumbs down”. E.g.

- Does everyone have their Colab environment ready?
- Are you clear on the key architectural differences between our three models?
- Do you feel confident about implementing SLMs in your own projects?



# Q&A





# Break

O'REILLY®